



# Computer vision in the new era of Artificial Intelligence and Deep Learning

## Visión por computador en la nueva era de la Inteligencia Artificial y el Deep Learning

**Rubén Usamentiaga\*, Alberto Fernández°**

**\* University of Oviedo**

**° TSK**

Gijón (Spain)  
5 – 16 April 2021



<https://github.com/albertofernandezvillan/computer-vision-and-deep-learning-course>

# TensorFlow and Keras

## Introduction to TensorFlow and Keras

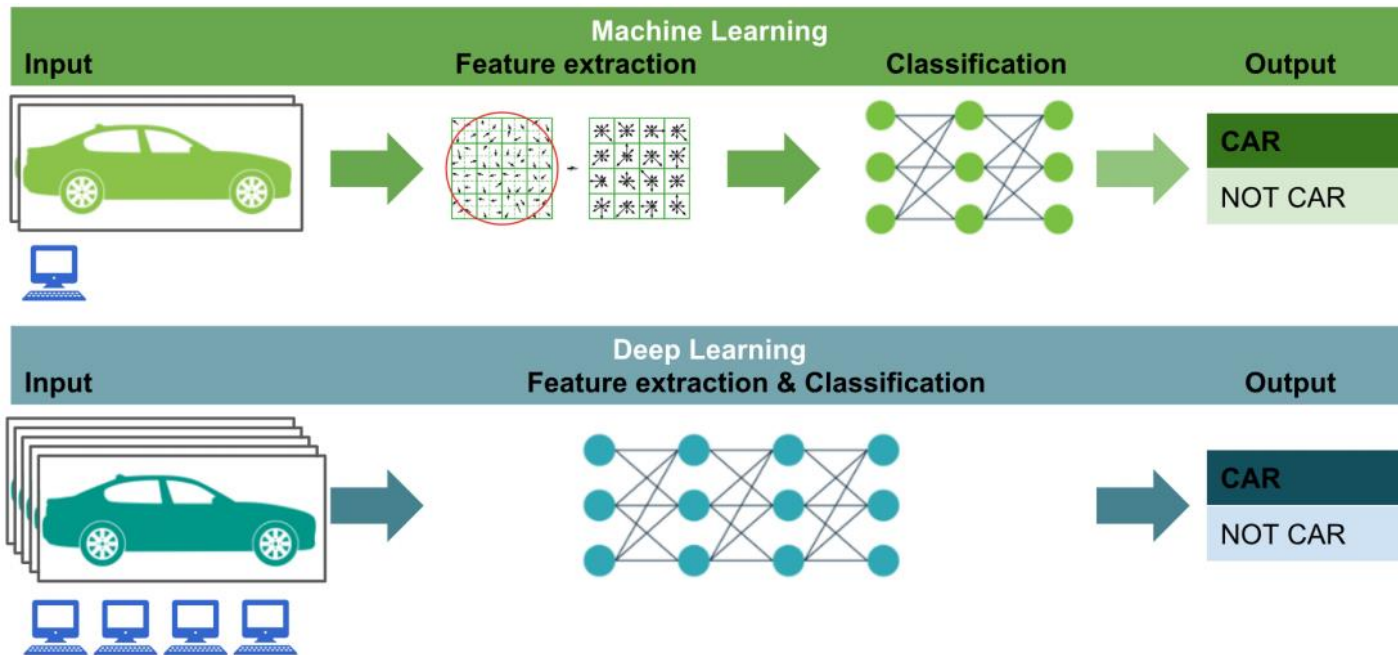




- [keras applications prediction.ipynb](#)
  - [set up kaggle api in colab.ipynb](#)
  - [keras applications feature extraction for classification.ipynb](#)
  - [keras applications feature extraction for clustering.ipynb](#)
  - [keras imagedatagenerator and dataset augmentation.ipynb](#)
  - [keras applications transfer learning.ipynb](#)
- 
- [keras applications prediction.ipynb](#)
  - [set up kaggle api in colab.ipynb](#)
  - [keras applications feature extraction for classification.ipynb](#)
  - [keras applications feature extraction for clustering.ipynb](#)
  - [keras imagedatagenerator and dataset augmentation.ipynb](#)
  - [keras applications transfer learning.ipynb](#)



<https://github.com/albertofernandezvillan/computer-vision-and-deep-learning-course>

# Machine learning vs deep learning



# Deep learning frameworks

Deep learning frameworks introduced in this course



Keras

 PyTorch

# Introduction to Keras

- Models API Different ways to create Keras models
- Layers API API for creating layers and also using many built-in layers
- Callbacks API Performing “actions” at various stages of training (e.g. At start or end of an epoch)
- Data preprocessing Transfor raw data on disk to a Dataset. Also other functions to work with images (e.g. load an image into PIL format)
- Optimizers Different available optimizers (e.g. SGD, RMSprop, Adam)
- Metrics Many available metrics to judge the permormance of your model
- Losses Many available losses, which compute the quantity that a model should seek to minimize during training
- Built-in datasets Few toy datasets (already-vectorized, in Numpy format) that can be used for debugging a model or creating simple code examples. See also [TensorFlow Datasets](#)
- Keras Applications Deep learning models that are made available alongside pre-trained weights. These models can be used for prediction, feature extraction, and fine-tuning
- Utilities Many utilities for: 1) model plotting, 2) serialization, 3) Python NumPy, ...

# Introduction to Keras Applications

**Keras Applications:** Keras Applications are deep learning models that are made available alongside pre-trained weights. These models can be used for 1) prediction, 2) feature extraction, and 3) fine-tuning.

## Example for prediction

```
from tensorflow.keras.applications.resnet50 import ResNet50
```

```
model = ResNet50(weights='imagenet')
```

Instantiates the ResNet50 architecture with the weights pre-trained on ImageNet (note that `include_top` is True by default)





# Introduction to Keras Applications

**Keras Applications:** Keras Applications are deep learning models that are made available alongside pre-trained weights. These models can be used for 1) prediction, 2) feature extraction, and 3) fine-tuning.

## Example for prediction

```
from tensorflow.keras.applications.resnet50 import ResNet50
from tensorflow.keras.preprocessing import image
from tensorflow.keras.applications.resnet50 import preprocess_input, decode_predictions
```

```
model = ResNet50(weights='imagenet')
```

Instantiates the ResNet50 architecture with the weights pre-trained on ImageNet (note that `include_top` is True by default)

```
img_loaded = image.load_img(IMG_NAME, target_size=(224, 224))
```

Load the image (PIL format)

```
x = image.img_to_array(img_loaded)
x = np.expand_dims(x, axis=0)
x = preprocess_input(x)
```

Converts a PIL Image instance to a NumPy array

Shape: (1, 224, 224, 3)

The images are converted from RGB to BGR, then each color channel is zero-centered

```
preds = model.predict(x)
```

Get the predictions

```
decoded_preds = decode_predictions(preds, top=3)[0]
```

decode the results into a list of tuples (class, description, probability)

# Introduction to Keras Applications

**Keras Applications:** Keras Applications are deep learning models that are made available alongside pre-trained weights. These models can be used for 1) prediction, 2) feature extraction, and 3) fine-tuning.

## Example for prediction

```
from tensorflow.keras.applications.resnet50 import ResNet50
from tensorflow.keras.preprocessing import image
from tensorflow.keras.applications.resnet50 import preprocess_input
```

```
model = ResNet50(weights='imagenet')
```

```
img_loaded = image.load_img(IMG_NAME, target_size=(224, 224))
```

```
x = image.img_to_array(img_loaded)
x = np.expand_dims(x, axis=0)
x = preprocess_input(x)
```

```
preds = model.predict(x)
```

Get the predictions

```
decoded_preds = decode_predictions(preds, top=3)[0]
```

decode the results into a list of tuples  
(class, description, probability)



tions

nts pre-  
y default

format)

y array


GR, then



# Introduction to Keras Applications

## Example for feature extraction in a classification problem

- Convolution layers extract features from the image
- Fully connected layers classify the image using extracted features.




`include_top` parameter when instanciating the model indicates if including (`True`) or not including (`False`) the fully connected layers at the top of the network

# Introduction to Keras Applications

## Example for feature extraction in a classification problem

VGG16

```
tf.keras.applications.VGG16(  
    include_top=True,  
    weights="imagenet",  
    input_tensor=None,  
    input_shape=None,  
    pooling=None,  
    classes=1000,  
    classifier_activation="softmax")
```



# Introduction to Keras Applications

VGG16

## Example for feature extraction in a classification problem

**We will be using the first one. See the notebook for more information**

```
model = VGG16(weights="imagenet", include_top=False, input_shape=(224, 224, 3))
• Output from block5_pool (MaxPooling2D): (None, 7, 7, 512)
```

```
model_2 = VGG16(weights="imagenet", include_top=False, input_shape=(112, 112, 3))
• Output from block5_pool (MaxPooling2D) (None, 3, 3, 512)
```

```
model_pooling = VGG16(weights="imagenet", include_top=False, input_shape=(224, 224, 3),
pooling='max')
• Output from global_max_pooling2d (GlobalMaxPooling2D) (None, 512)
```

```
model_2_pooling = VGG16(weights="imagenet", include_top=False, input_shape=(112, 112, 3),
pooling='max')
• Output from global_max_pooling2d_1 (GlobalMaxPooling2D) (None, 512)
```

```
model_flatten = tf.keras.models.Sequential()
model_flatten.add(VGG16(weights="imagenet", include_top=False), input_shape=(224, 224, 3))
model_flatten.add(tf.keras.layers.Flatten())
• Output from flatten: (None, 25088)
```

# Introduction to Keras Applications

## Example for feature extraction in a classification problem

VGG16

```
model = VGG16(weights="imagenet", include_top=False, input_shape=(224, 224, 3))
```

- Output from block5\_pool (MaxPooling2D): **(None, 7, 7, 512)**

```
import numpy as np

def get_features(path_image):
    img = load_img(path_image, target_size=(224, 224))
    img = img_to_array(img)
    img = np.expand_dims(img, axis=0)
    img = preprocess_input(img)

    vgg16_feats= model.predict(img)
    vgg16_feats_flat = vgg16_feats.flatten()

    return vgg16_feats_flat
```

For each image, this function returns a flatten array with a shape of (25088,)

# Introduction to Keras Applications

## Example for feature extraction in a classification problem



N\_DOGS = 100



N\_CATS = 100

```
vgg16_feature_list = []  
labels = []
```

```
for each image:  
    features = get_features(path_image)  
    vgg16_feature_list.append(features)  
    labels.append(class_name)
```

```
vgg16_feature_list_np = np.array(vgg16_feature_list) #(200,25088)  
labels_np = np.array(labels) #(200,)
```

```
from sklearn.linear_model import LogisticRegression
```


```
log_reg_model = LogisticRegression()  
log_reg_model.fit(vgg16_feature_list_np, labels_np)
```

```
label =  
log_reg_model.predict(get_features(path_image).reshape(1,-1))[0]
```



# Introduction to Keras Applications

## Example for feature extraction in a clustering problem



N\_DOGS = 100



N\_CATS = 100

```
vgg16_feature_list = []  
labels = []
```

```
for each image:  
    features = get_features(path_image)  
    vgg16_feature_list.append(features)  
    labels.append(class_name)
```

```
vgg16_feature_list_np = np.array(vgg16_feature_list) #(200,25088)  
labels_np = np.array(labels) #(200,)
```

```
from sklearn.cluster import KMeans
```

```
kmeans = KMeans(n_clusters=2, random_state=0)  
Kmeans. fit(vgg16_feature_list_np)
```

```
label =  
log_reg_model.predict(get_features(path_image).reshape(1,-1))[0]
```





# Introduction to Keras Applications

## Example for feature extraction in a clustering problem



N\_DOGS = 100



N\_CATS = 100

```
vgg16_feature_list = []  
labels = []
```

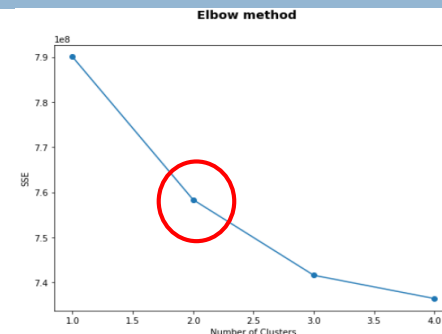
```
for each image:  
    features = get_features(path_image)  
    vgg16_feature_list.append(features)  
    labels.append(class_name)
```

```
vgg16_feature_list_np = np.array(vgg16_feature_list) # (200, 25088)  
labels_np = np.array(labels) # (200,)
```

```
from sklearn.cluster import KMeans
```

```
kmeans = KMeans(n_clusters=2, random_state=0)  
kmeans.fit(vgg16_feature_list_np)
```

```
label =  
log_reg_model.predict(get_features(path_image).reshape(1,-1))[0]
```



elbow method (see notebook)



# Data augmentation

[https://www.tensorflow.org/api\\_docs/python/tf/keras/preprocessing/image/ImageDataGenerator](https://www.tensorflow.org/api_docs/python/tf/keras/preprocessing/image/ImageDataGenerator)

```
import tensorflow as tf
```

```
data_generator =
```

```
tf.keras.preprocessing.image.ImageDataGenerator(rotation_range=30)
```

```
image_iterator = data_generator.flow(images)
```

```
show_augs(image_iterator)
```



# Data augmentation

[https://www.tensorflow.org/api\\_docs/python/tf/keras/preprocessing/image/ImageDataGenerator](https://www.tensorflow.org/api_docs/python/tf/keras/preprocessing/image/ImageDataGenerator)

```
import tensorflow as tf
```

```
data_generator = tf.keras.preprocessing.image.ImageDataGenerator(zoom_range=0.25)
```

```
image_iterator = data_generator.flow(images)
```

```
show_augs(image_iterator)
```





# Data augmentation

[https://www.tensorflow.org/api\\_docs/python/tf/keras/preprocessing/image/ImageDataGenerator](https://www.tensorflow.org/api_docs/python/tf/keras/preprocessing/image/ImageDataGenerator)

```
import tensorflow as tf
```

```
data_generator = tf.keras.preprocessing.image.ImageDataGenerator(width_shift_range=0.3)
```

```
image_iterator = data_generator.flow(images)
```

```
show_augs(image_iterator)
```



# Data augmentation

```
data_generator = tf.keras.preprocessing.image.ImageDataGenerator(  
    rescale=1. / 255,  
    zoom_range=0.2,  
    rotation_range = 5,  
    horizontal_flip=True)
```




# Introduction to Keras Applications

## Example for transfer learning

Image taken from [here](#)

### Multiclass classification

#### Pre-training



convolutional layers



dense layers

64% tabby cat

33% Siamese cat


⋮

0.1% wooden spoon


⋮

### Binary classification

#### Transfer learning



convolutional layers (frozen)



new dense layers

92% Alien

8% Predator



# TensorFlow and Keras

Introduction to TensorFlow and Keras



<https://github.com/albertofernandezvillan/computer-vision-and-deep-learning-course>