



# Computer vision in the new era of Artificial Intelligence and Deep Learning

## Visión por computador en la nueva era de la Inteligencia Artificial y el Deep Learning

**Rubén Usamentiaga\*, Alberto Fernández°**

**\* University of Oviedo**

**° TSK**

Gijón (Spain)  
5 – 16 April 2021



<https://github.com/albertofernandezvillan/computer-vision-and-deep-learning-course>

# TensorFlow and Keras

## Introduction to TensorFlow and Keras

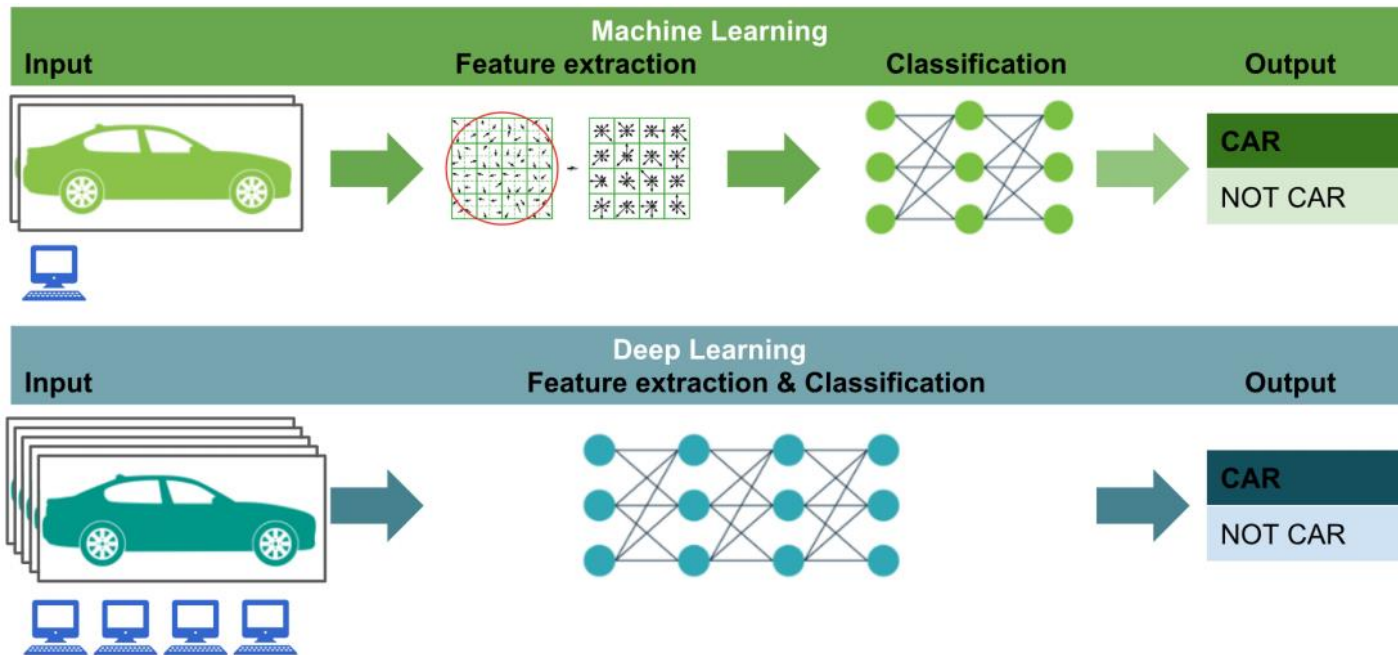




- [keras applications prediction.ipynb](#)
  - [set up kaggle api in colab.ipynb](#)
  - [keras applications feature extraction for classification.ipynb](#)
  - [keras applications feature extraction for clustering.ipynb](#)
  - [keras imagedatagenerator and dataset augmentation.ipynb](#)
  - [keras applications transfer learning dogs.ipynb](#)
- 
- [keras applications prediction.ipynb](#)
  - [set up kaggle api in colab.ipynb](#)
  - [keras applications feature extraction for classification.ipynb](#)
  - [keras applications feature extraction for clustering.ipynb](#)
  - [keras imagedatagenerator and dataset augmentation.ipynb](#)
  - [keras applications transfer learning dogs.ipynb](#)



<https://github.com/albertofernandezvillan/computer-vision-and-deep-learning-course>

# Machine learning vs deep learning



# Deep learning frameworks

Deep learning frameworks introduced in this course



Keras

 PyTorch

# Introduction to Keras

- Models API Different ways to create Keras models
- Layers API API for creating layers and also using many built-in layers
- Callbacks API Performing “actions” at various stages of training (e.g. At start or end of an epoch)
- Data preprocessing Transfor raw data on disk to a Dataset. Also other functions to work with images (e.g. load an image into PIL format)
- Optimizers Different available optimizers (e.g. SGD, RMSprop, Adam)
- Metrics Many available metrics to judge the permormance of your model
- Losses Many available losses, which compute the quantity that a model should seek to minimize during training
- Built-in datasets Few toy datasets (already-vectorized, in Numpy format) that can be used for debugging a model or creating simple code examples. See also [TensorFlow Datasets](#)
- Keras Applications Deep learning models that are made available alongside pre-trained weights. These models can be used for prediction, feature extraction, and fine-tuning
- Utilities Many utilities for: 1) model plotting, 2) serialization, 3) Python, ...

# Introduction to Keras Applications

**Keras Applications:** Keras Applications are deep learning models that are made available alongside pre-trained weights. These models can be used for 1) prediction, 2) feature extraction, and 3) fine-tuning.

## Example for prediction

```
from tensorflow.keras.applications.resnet50 import ResNet50
```

```
model = ResNet50(weights='imagenet')
```

Instantiates the ResNet50 architecture with the weights pre-trained on ImageNet (note that `include_top` is True by default)





# Introduction to Keras Applications

## Example for prediction

```
from tensorflow.keras.applications.resnet50 import ResNet50
from tensorflow.keras.preprocessing import image
from tensorflow.keras.applications.resnet50 import preprocess_input, decode_predictions
```

```
model = ResNet50(weights='imagenet')
```

Instantiates the ResNet50 architecture with the weights pre-trained on ImageNet (note that `include_top` is True by default)

```
img_loaded = image.load_img(IMG_NAME, target_size=(224, 224))
```

Load the image (PIL format)

```
x = image.img_to_array(img_loaded)
x = np.expand_dims(x, axis=0)
x = preprocess_input(x)
```

Converts a PIL Image instance to a NumPy array

Shape: (1, 224, 224, 3)

The images are converted from RGB to BGR, then each color channel is zero-centered

```
preds = model.predict(x)
```

Get the predictions

```
decoded_preds = decode_predictions(preds, top=3)[0]
```

decode the results into a list of tuples (class, description, probability)

# Introduction to Keras Applications

## Example for prediction

```
from tensorflow.keras.applications.resnet50 import ResNet50
from tensorflow.keras.preprocessing import image
from tensorflow.keras.applications.resnet50 import preprocess_input, decode_predictions
```

```
model = ResNet50(weights='imagenet')
```

```
img_loaded = image.load_img(IMG_NAME, target_size=(224, 224))
```

```
x = image.img_to_array(img_loaded)
x = np.expand_dims(x, axis=0)
x = preprocess_input(x)
```

```
preds = model.predict(x)
```

```
decoded_preds = decode_predictions(preds, top=3)[0]
```

Inst  
tra

Get the predictions



tions

nts pre-  
y default

format)

y array

GR, then


decode the results into a list of tuples  
(class, description, probability)



# Introduction to Keras Applications

## Example for feature extraction in a classification problem

- Convolution layers extract features from the image
- Fully connected layers classify the image using extracted features.




`include_top` parameter when instanciating the model indicates if including (`True`) or not including (`False`) the fully connected layers at the top of the network

# Introduction to Keras Applications

## Example for feature extraction in a classification problem

VGG16

```
tf.keras.applications.VGG16(  
    include_top=True,  
    weights="imagenet",  
    input_tensor=None,  
    input_shape=None,  
    pooling=None,  
    classes=1000,  
    classifier_activation="softmax")
```



# Introduction to Keras Applications

## Example for feature extraction in a classification problem

VGG16

**We will use the first one. See the notebook for more information**

```
model = VGG16(weights="imagenet", include_top=False, input_shape=(224, 224, 3))  
• Output from block5_pool (MaxPooling2D): (None, 7, 7, 512)
```

```
model_2 = VGG16(weights="imagenet", include_top=False, input_shape=(112, 112, 3))  
• Output from block5_pool (MaxPooling2D) (None, 3, 3, 512)
```

```
model_pooling = VGG16(weights="imagenet", include_top=False, input_shape=(224, 224, 3),  
pooling='max')  
• Output from global_max_pooling2d (GlobalMaxPooling2D) (None, 512)
```

```
model_2_pooling = VGG16(weights="imagenet", include_top=False, input_shape=(112, 112, 3),  
pooling='max')  
• Output from global_max_pooling2d_1 (GlobalMaxPooling2D) (None, 512)
```

```
model_flatten = tf.keras.models.Sequential()  
model_flatten.add(VGG16(weights="imagenet", include_top=False), input_shape=(224, 224, 3))  
model_flatten.add(tf.keras.layers.Flatten())  
• Output from flatten: (None, 25088)
```

# Introduction to Keras Applications

## Example for feature extraction in a classification problem

VGG16

```
model = VGG16(weights="imagenet", include_top=False, input_shape=(224, 224, 3))
```

- Output from block5\_pool (MaxPooling2D): **(None, 7, 7, 512)**

```
import numpy as np

def get_features(path_image):
    img = load_img(path_image, target_size=(224, 224))
    img = img_to_array(img)
    img = np.expand_dims(img, axis=0)
    img = preprocess_input(img)

    vgg16_feats= model.predict(img)
    vgg16_feats_flat = vgg16_feats.flatten()

    return vgg16_feats_flat
```

For each image, this function returns a flatten array with a shape of (25088,)

# Introduction to Keras Applications

## Example for feature extraction in a classification problem



N\_DOGS = 100



N\_CATS = 100

```
vgg16_feature_list = []  
labels = []
```

```
for each image:  
    features = get_features(path_image)  
    vgg16_feature_list.append(features)  
    labels.append(class_name)
```

```
vgg16_feature_list_np = np.array(vgg16_feature_list) #(200,25088)  
labels_np = np.array(labels) #(200,)
```

```
from sklearn.linear_model import LogisticRegression
```


```
log_reg_model = LogisticRegression()  
log_reg_model.fit(vgg16_feature_list_np, labels_np)
```

```
label =  
log_reg_model.predict(get_features(path_image).reshape(1,-1))[0]
```



# Introduction to Keras Applications

## Example for feature extraction in a clustering problem



N\_DOGS = 100



N\_CATS = 100

```
vgg16_feature_list = []  
labels = []
```


```
for each image:  
    features = get_features(path_image)  
    vgg16_feature_list.append(features)  
    labels.append(class_name)
```

```
vgg16_feature_list_np = np.array(vgg16_feature_list) #(200,25088)  
labels_np = np.array(labels) #(200,)
```

```
from sklearn.cluster import KMeans
```

```
kmeans = KMeans(n_clusters=2, random_state=0)  
Kmeans.fit(vgg16_feature_list_np)
```

```
label =  
log_reg_model.predict(get_features(path_image).reshape(1,-1))[0]
```





# Introduction to Keras Applications

## Example for feature extraction in a clustering problem



N\_DOGS = 100



N\_CATS = 100

```
vgg16_feature_list = []  
labels = []
```


```
for each image:  
    features = get_features(path_image)  
    vgg16_feature_list.append(features)  
    labels.append(class_name)
```

```
vgg16_feature_list_np = np.array(vgg16_feature_list) # (200, 25088)  
labels_np = np.array(labels) # (200,)
```

```
from sklearn.cluster import KMeans
```

```
kmeans = KMeans(n_clusters=2, random_state=0)  
kmeans.fit(vgg16_feature_list_np)
```

```
label =  
log_reg_model.predict(get_features(path_image).reshape(1,-1))[0]
```



elbow method (see notebook)



# Data augmentation

[https://www.tensorflow.org/api\\_docs/python/tf/keras/preprocessing/image/ImageDataGenerator](https://www.tensorflow.org/api_docs/python/tf/keras/preprocessing/image/ImageDataGenerator)

```
import tensorflow as tf
```

```
data_generator =
```

```
tf.keras.preprocessing.image.ImageDataGenerator(rotation_range=30)
```

```
image_iterator = data_generator.flow(images)
```

```
show_augs(image_iterator)
```



# Data augmentation

[https://www.tensorflow.org/api\\_docs/python/tf/keras/preprocessing/image/ImageDataGenerator](https://www.tensorflow.org/api_docs/python/tf/keras/preprocessing/image/ImageDataGenerator)

```
import tensorflow as tf
```

```
data_generator = tf.keras.preprocessing.image.ImageDataGenerator(zoom_range=0.25)
```

```
image_iterator = data_generator.flow(images)
```

```
show_augs(image_iterator)
```





# Data augmentation

[https://www.tensorflow.org/api\\_docs/python/tf/keras/preprocessing/image/ImageDataGenerator](https://www.tensorflow.org/api_docs/python/tf/keras/preprocessing/image/ImageDataGenerator)

```
import tensorflow as tf
```

```
data_generator = tf.keras.preprocessing.image.ImageDataGenerator(width_shift_range=0.3)
```

```
image_iterator = data_generator.flow(images)
```

```
show_augs(image_iterator)
```



# Data augmentation

```
data_generator = tf.keras.preprocessing.image.ImageDataGenerator(  
    rescale=1. / 255,  
    zoom_range=0.2,  
    rotation_range = 5,  
    horizontal_flip=True)
```




# Introduction to Keras Applications

## Example for transfer learning

Image taken from [here](#)

### Multiclass classification



convolutional layers



dense layers

64% tabby cat


33% Siamese cat

⋮


0.1% wooden spoon

⋮

### Binary classification



convolutional layers (frozen)



new dense layers

92% Alien


8% Predator



# Introduction to Keras Applications

Example for transfer learning

See the notebook for more information




Base model	resolution
EfficientNetB0	224
EfficientNetB1	240
EfficientNetB2	260
EfficientNetB3	300
EfficientNetB4	380
EfficientNetB5	456
EfficientNetB6	528
EfficientNetB7	600

# Introduction to Keras Applications

```
from tensorflow.keras.applications import EfficientNetB0
```

```
base_model_eff = EfficientNetB0(include_top=False,  
input_shape=(INPUT_SHAPE[0], INPUT_SHAPE[1], 3),  
weights="imagenet")
```



output is a 4D tensor

(None, 7, 7, 1280)

```
base_model_eff = EfficientNetB0(include_top=False, input_shape=(INPUT_SHAPE[0],  
INPUT_SHAPE[1], 3), weights="imagenet")  
model = models.Sequential()  
model.add(base_model_eff)  
model.add(layers.GlobalMaxPooling2D())  
model.summary()
```

**GlobalMaxPooling2D() results in a much smaller number of features compared to the Flatten() layer**

Model: "sequential"



Layer (type)	Output Shape	Param #
efficientnetb0 (Functional)	(None, 7, 7, 1280)	4049571
global_max_pooling2d (Global)	(None, 1280)	0
Total params: 4,049,571		
Trainable params: 4,007,548		
Non-trainable params: 42,023		

output is a 2D tensor

(None, 1280)

[tf.keras.layers.GlobalMaxPool2D\(\)](#)

```
#(batch_size, rows, cols, channels)  
input_shape = (1, 3, 3, 3)  
  
x = tf.random.normal(input_shape)  
y = tf.keras.layers.GlobalMaxPool2D()(x)  
  
#(batch_size, channels)  
print(y.shape) # (1, 3)
```



**See the notebook for more information**

# TensorFlow and Keras

Introduction to TensorFlow and Keras



<https://github.com/albertofernandezvillan/computer-vision-and-deep-learning-course>